

MACHINE LEARNING

ASSIGNMENT NO.-01

Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

Perform following tasks:

1. Pre-process the dataset.

2. Identify outliers.

3. Check the correlation.

4. Implement linear regression and random forest regression models.

5. Evaluate the models and compare their respective scores like R2, RMSE, etc.

Dataset link: <https://www.kaggle.com/datasets/yasserh/uber-fares-datase>

Import necessary libraries

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import r2_score, mean_squared_error

Load the dataset

data = pd.read_csv("Uber.csv")

data

Unnamed: 0		key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.723217	1
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.750325	1
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.772647	1
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.803349	3
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.761247	5
...
199995	42598914	2012-10-28 10:49:00.000000053	3.0	2012-10-28 10:49:00 UTC	-73.987042	40.739367	-73.986525	40.740297	1
199996	16382965	2014-03-14 01:09:00.00000008	7.5	2014-03-14 01:09:00 UTC	-73.984722	40.736837	-74.006672	40.739620	1
199997	27804658	2009-06-29 00:42:00.000000078	30.9	2009-06-29 00:42:00 UTC	-73.986017	40.756487	-73.858957	40.692588	2
199998	20259894	2015-05-20 14:56:25.00000004	14.5	2015-05-20 14:56:25 UTC	-73.997124	40.725452	-73.983215	40.695415	1
199999	11951496	2010-05-15 04:08:00.000000076	14.1	2010-05-15 04:08:00 UTC	-73.984395	40.720077	-73.985508	40.768793	1

200000 rows × 9 columns

1. Pre-process the dataset

Remove unnecessary column

```
data["pickup_datetime"] = pd.to_datetime(data["pickup_datetime"])
```

```
missing_values = data.isnull().sum()
```

```
print("Missing values in the dataset:")
```

```
print(missing_values)
```

Handle missing values

We can choose to drop rows with missing values or fill them with appropriate values.

```
data.dropna(inplace=True)
```

To fill missing values with the mean value of the column:

```
# data.fillna(data.mean(), inplace=True)
```

Ensure there are no more missing values

```
missing_values = data.isnull().sum()
```

```
print("Missing values after handling:")
```

```
print(missing_values)
```

2. Identify outliers

visualization to detect outliers.

```
sns.boxplot(x=data["fare_amount"])
```

```
plt.show()
```

Missing values in the dataset:


key	0
fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	1
dropoff_latitude	1
passenger_count	0

dtype: int64

Missing values after handling:

key	0
fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	0
dropoff_latitude	0
passenger_count	0

dtype: int64



```
# Calculate the IQR for the 'fare_amount' column
```

```
Q1 = data["fare_amount"].quantile(0.25)
```

```
Q3 = data["fare_amount"].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
# Define a threshold (e.g., 1.5 times the IQR) to identify outliers
```

```
threshold = 1.5
```


```
lower_bound = Q1 - threshold * IQR
```

```
upper_bound = Q3 + threshold * IQR
```

```
# Remove outliers
```



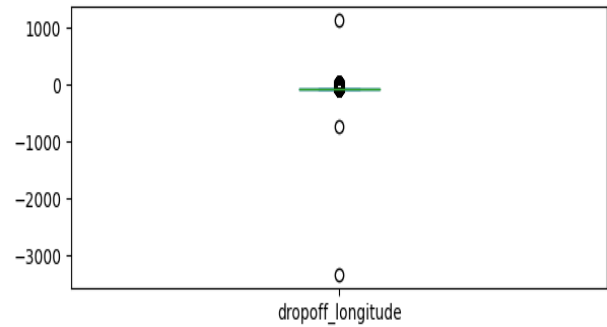


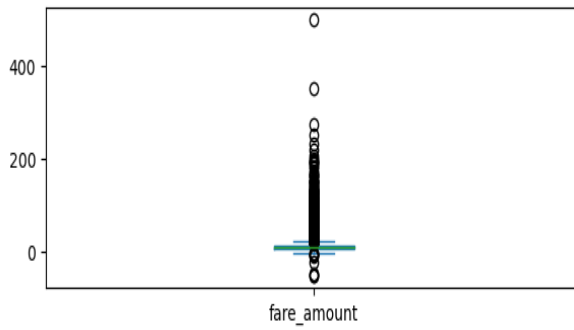
```
data_no_outliers = data[(data["fare_amount"] >= lower_bound) & (data["fare_amount"] <= upper_bound)]
```

```
# Visualize the 'fare_amount' distribution without outliers
sns.boxplot(x=data_no_outliers["fare_amount"])
plt.show()
```

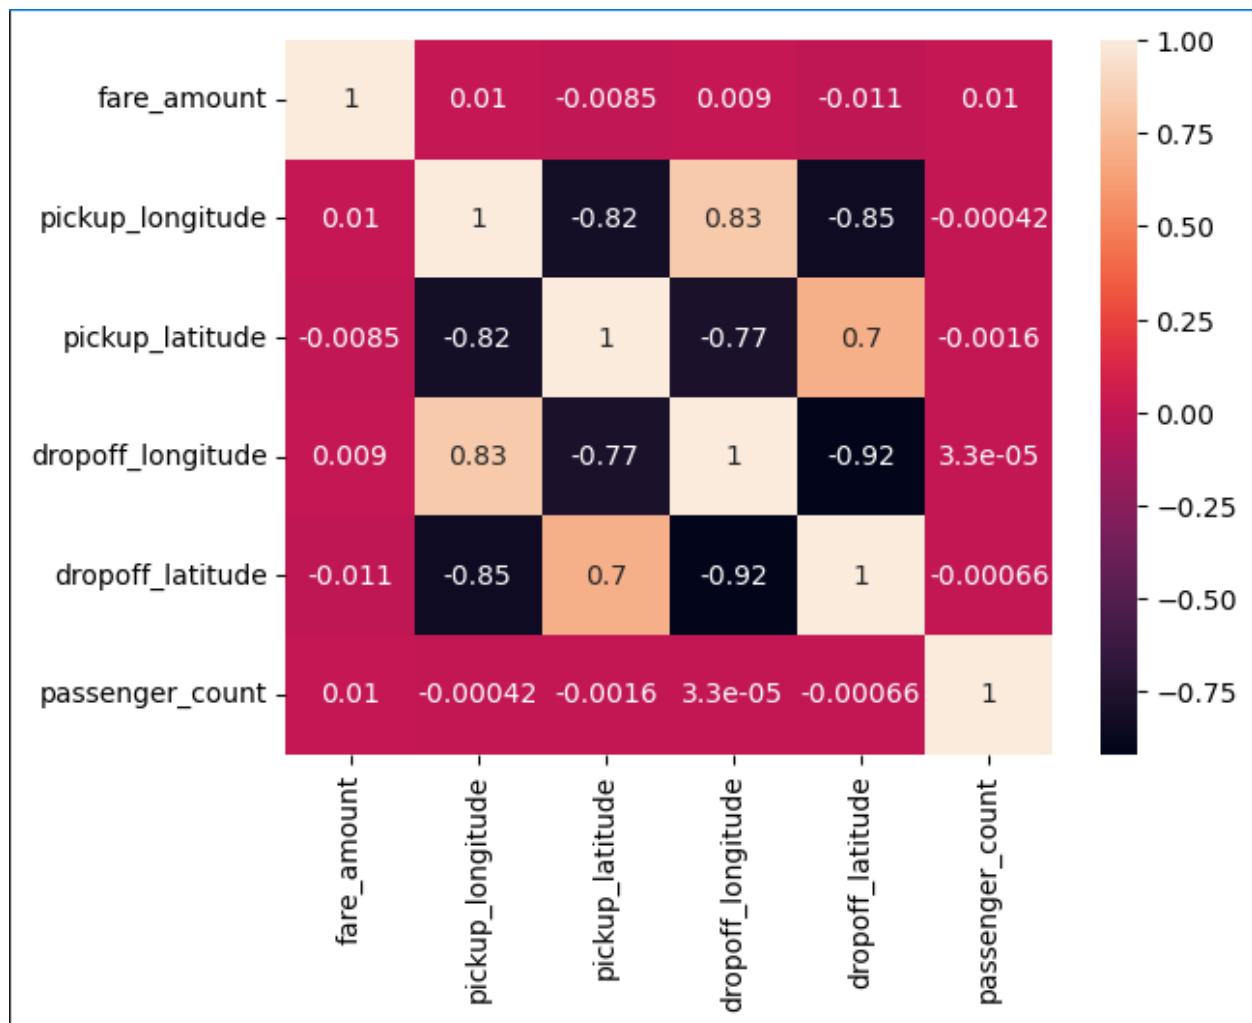


```
data.plot(kind="box",subplots=True, layout=(7, 2), figsize=(15, 20))
```

```
fare_amount      AxesSubplot(0.125,0.786098;0.352273x0.0939024)
pickup_longitude AxesSubplot(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude  AxesSubplot(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude AxesSubplot(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude AxesSubplot(0.125,0.560732;0.352273x0.0939024)
passenger_count  AxesSubplot(0.547727,0.560732;0.352273x0.0939024)
dtype: object
```



```
# 3. Check the correlation
# Determine the correlation between features and the target variable
(fare_amount).
correlation_matrix = data.corr()
sns.heatmap(correlation_matrix, annot=True)
plt.show()
```



```
# 4. Implement linear regression and random forest regression models
# Split the data into features and target variable
X = data[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
'dropoff_latitude', 'passenger_count']]
y = data['fare_amount'] #Target
y
```

```
0      7.5
1      7.7
2     12.9
3      5.3
4     16.0
...
199995   3.0
199996   7.5
199997  30.9
199998  14.5
199999  14.1
Name: fare_amount, Length: 199999, dtype: float64
```

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
# Create and train the linear regression model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```

```
# Create and train the random forest regression model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
```

```
# 5. Evaluate the models
# Predict the values
y_pred_lr = lr_model.predict(X_test)
y_pred_lr
print("Linear Model:", y_pred_lr)
y_pred_rf = rf_model.predict(X_test)
print("Random Forest Model:", y_pred_rf)
```

```
# Calculate R-squared (R2) and Root Mean Squared Error (RMSE) for both models
r2_lr = r2_score(y_test, y_pred_lr)
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))
```

```
# Compare the scores
print("Linear Regression - R2:", r2_lr)
print("Linear Regression - RMSE:", rmse_lr)
```

```
Linear Regression - R2: 0.00034152697863043535
Linear Regression - RMSE: 10.197470623964248
```

```
r2_rf = r2_score(y_test, y_pred_rf)
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))

print("Random Forest Regression R2:", r2_rf)
print("Random Forest Regression RMSE:", rmse_rf)
```

```
Random Forest Regression R2: 0.7011790407391916
Random Forest Regression RMSE: 5.575350372469675
```